

Time series decomposition

[Print-friendly PDF version](#)

Before you start

This is an in-class lab: work through the material below at your own pace (in pairs or small groups is fine) — I'll walk around to help, and we'll stop along the way to discuss as a group.

This builds on the moving averages, classical decomposition, X11/SEATS and STL material from the [decomposition chapter](#). Today you apply it: decompose real series into trend-cycle, seasonal and remainder components, compare methods, and interpret what you get.

```
library(fpp3)
library(tidyverse)
```

Cheat sheet: decomposition methods

Method	Syntax
Classical	<code>model(classical_decomposition(variable, type = "additive"))</code>
X11	<code>model(X_13ARIMA_SEATS(variable ~ x11()))</code>
SEATS	<code>model(X_13ARIMA_SEATS(variable ~ seats()))</code>
STL	<code>model(STL(variable ~ trend(window = ...) + season(window = ...), robust = ...))</code>

After fitting, use `components()` to extract the trend-cycle/seasonal/remainder columns, then `autoplot()` to plot them.

Task 1: Offshore wind power

Use the `windpower` data from `ban430data` (the same one from the [Time series graphics lab](#)) again — this time for decomposition.

```
# install.packages("remotes")
remotes::install_github("holleland/ban430data")
library(ban430data)
data(windpower)
```

- Create a tsibble containing the **daily** wind power data from **Sørlig Nordsjø 2** only (aggregate **windpower** from hourly to daily with the mean, same as in the graphics lab).
- Decompose the series into trend-cycle T_t , season S_t and remainder R_t , using a suitable decomposition method. Which method(s) from the cheat sheet above are *not* usable here, and why? Why did you choose the method you did?
- Is there a visible seasonal pattern? A trend? How does this compare to what you found with `gg_season()`/`gg_subseries()` in the graphics lab?

Task 2: Norwegian wholesale and retail sales index

The data is a monthly index for *Retail trade, except of motor vehicles and motorcycles*, from Statistics Norway (Jan 2000 onwards).

```
wholesale <- read.csv2(
  "https://raw.githubusercontent.com/holleland/BAN430/master/data/wholesale_and_retails_index.csv",
  sep = ";")
head(wholesale, 3)
```

	month	wholesale_and_retail_sales_index
1	2000M01	50.5
2	2000M02	49.3
3	2000M03	53.8

- Convert the **month** column to a **yearmonth** type and turn **wholesale** into a tsibble.
- Make a time plot.
- Decompose the series using the classical, X11, SEATS and STL methods. Can you detect any prominent differences between the methods?
- Try adjusting the trend and season windows of the STL (default values are 21 and 11 respectively). What happens?
- Using your method of choice, plot the seasonally-adjusted time series.
- Using your method of choice, plot the detrended series.
- Optional:* Implement your own additive classical decomposition by hand (moving averages + seasonal averaging), and compare to `classical_decomposition()`. (See [Chapter 3 exercises](#), [exercise 6](#) if you want to check your approach against a worked solution.)

Wrap-up discussion

For both series today: did the different decomposition methods roughly agree, or did they tell noticeably different stories? When might you prefer STL over classical decomposition, or vice versa? What made X11/SEATS unsuitable for the daily wind power data?